SAME CRASH, DIFFERENT STORY

Can wording change an AI's moral decisions?

Explainable AI course project

---

1 / SAME CONSEQUENCES

A car cannot stop. Its options can be described through deaths or survivors. The facts stay the same. Does the decision?

One uncertain example

STAY: 1 death, certainly.
SWERVE: 20% chance of 5 deaths; 80% chance of no deaths.
Ten people in total; all other people survive.

Both actions have one expected death, but different risks. We do not prescribe a morally correct answer.

2 / CONTROLLED TEST

8 scenarios
3 equivalent descriptions
2 instruction conditions
5 repeats per prompt
= 240 planned responses

Four families: clear, close, equal expected harm, and lower expected harm with risk. Each is mirrored.

Compare repeated identical prompts against equivalent descriptions.

Intervention: restate both options symmetrically before choosing.

3 / WHAT WE FOUND

240/240 valid responses
Model: gpt-6.1-sol

Standard

Same-prompt disagreement: 5.8%
Cross-description disagreement: 5.0%
Excess: -0.8 percentage points
8/8 complete scenarios

Symmetric restatement

Same-prompt disagreement: 1.7%
Cross-description disagreement: 1.7%
Excess: 0.0 percentage points
8/8 complete scenarios

Variation appeared without excess framing disagreement.

We test behavioral consistency, not moral correctness or driving safety. One model alias, eight designed scenarios and five repeats limit generalization. Pairwise comparisons share responses; no population significance claim is made.

Next: independent prompt-equivalence and explanation review; more scenarios; test whether explanations help people predict decisions.

Explore: predict a choice, change the wording, reveal the recorded decisions and explanations.

Grounding: Tversky & Kahneman, Science (1981); Awad et al., Nature (2018); Turpin et al., NeurIPS (2023).